

Wisconsin Diagnostic Breast Cancer (WDBC)

Marina Vallejo Vallés (marina.vallejo01@estudiant.upf.edu (mailto:marina.vallejo01@estudiant.upf.edu))

Last update: 10 abril, 2022

- EXPLORATORY DATA ANALYSIS
- MODEL BUILDING
 - KNN
 - Decision Tree
 - Neural Network

Brief introduction.

The Wisconsin Diagnostic Breast Cancer (WDBC) contains features computed from a digitized image of a fine needle aspirate (FNA) of a breast mass.

In the literature they report that the two diagnosis results: B = benign, M = malignant are linearly separable using all 30 input features available in the data set. They also report that the best predictive accuracy is obtained using one separating plane in the 3-D space of Worst Area, Worst Smoothness and Mean Texture. With an estimated accuracy 97.5% using repeated 10-fold cross-validations.

EXPLORATORY DATA ANALYSIS

Work with the data set.

```
# Access the data:
breast_dataset <- read.csv("data_WDBC.csv", sep=",")
# Summary of the data:
summary(breast_dataset)

##      id      diagnosis      radius_mean      texture_mean
## Min.   :      8670      Length:569      Min.    : 6.981      Min.    : 9.71
## 1st Qu.:    869218      Class :character      1st Qu.:11.700      1st Qu.:16.17
## Median :    906024      Mode  :character      Median :13.370      Median :18.84
## Mean   :    30371831      Mean    :14.127      Mean    :19.29
## 3rd Qu.:    8813129      3rd Qu.:15.780      3rd Qu.:21.80
## Max.   :    911320502      Max.    :28.110      Max.    :39.28
## perimeter_mean      area_mean      smoothness_mean      compactness_mean
## Min.   : 43.79      Min.    : 143.5      Min.    :0.05263      Min.    :0.01938
## 1st Qu.: 75.17      1st Qu.: 420.3      1st Qu.:0.08637      1st Qu.:0.06492
## Median : 86.24      Median : 551.1      Median :0.09587      Median :0.09263
## Mean   : 91.97      Mean    : 654.9      Mean    :0.09636      Mean    :0.10434
## 3rd Qu.:104.10      3rd Qu.: 782.7      3rd Qu.:0.10530      3rd Qu.:0.13040
## Max.   :188.50      Max.    :2501.0      Max.    :0.16340      Max.    :0.34540
## concavity_mean      concave.points_mean      symmetry_mean      fractal_dimension_mean
## Min.   :0.00000      Min.    :0.00000      Min.    :0.1060      Min.    :0.04996
## 1st Qu.:0.02956      1st Qu.:0.02031      1st Qu.:0.1619      1st Qu.:0.05770
## Median :0.06154      Median :0.03350      Median :0.1792      Median :0.06154
## Mean   :0.08880      Mean    :0.04892      Mean    :0.1812      Mean    :0.06280
## 3rd Qu.:0.13070      3rd Qu.:0.07400      3rd Qu.:0.1957      3rd Qu.:0.06612
## Max.   :0.42680      Max.    :0.20120      Max.    :0.3040      Max.    :0.09744
## radius_se      texture_se      perimeter_se      area_se
## Min.   :0.1115      Min.    :0.3602      Min.    : 0.757      Min.    : 6.802
## 1st Qu.:0.2324      1st Qu.:0.8339      1st Qu.: 1.606      1st Qu.:17.850
## Median :0.3242      Median :1.1080      Median : 2.287      Median :24.530
## Mean   :0.4052      Mean    :1.2169      Mean    : 2.866      Mean    :40.337
## 3rd Qu.:0.4789      3rd Qu.:1.4740      3rd Qu.: 3.357      3rd Qu.:45.190
## Max.   :2.8730      Max.    :4.8850      Max.    :21.980      Max.   :542.200
## smoothness_se      compactness_se      concavity_se      concave.points_se
## Min.   :0.001713      Min.    :0.002252      Min.    :0.00000      Min.    :0.000000
## 1st Qu.:0.005169      1st Qu.:0.013080      1st Qu.:0.01509      1st Qu.:0.007638
## Median :0.006380      Median :0.020450      Median :0.02589      Median :0.010930
## Mean   :0.007041      Mean    :0.025478      Mean    :0.03189      Mean    :0.011796
## 3rd Qu.:0.008146      3rd Qu.:0.032450      3rd Qu.:0.04205      3rd Qu.:0.014710
## Max.   :0.031130      Max.    :0.135400      Max.    :0.39600      Max.    :0.052790
## symmetry_se      fractal_dimension_se      radius_worst      texture_worst
## Min.   :0.007882      Min.    :0.0008948      Min.    : 7.93      Min.    :12.02
## 1st Qu.:0.015160      1st Qu.:0.0022480      1st Qu.:13.01      1st Qu.:21.08
## Median :0.018730      Median :0.0031870      Median :14.97      Median :25.41
## Mean   :0.020542      Mean    :0.0037949      Mean    :16.27      Mean    :25.68
## 3rd Qu.:0.023480      3rd Qu.:0.0045580      3rd Qu.:18.79      3rd Qu.:29.72
## Max.   :0.078950      Max.    :0.0298400      Max.    :36.04      Max.    :49.54
## perimeter_worst      area_worst      smoothness_worst      compactness_worst
## Min.   : 50.41      Min.    : 185.2      Min.    :0.07117      Min.    :0.02729
## 1st Qu.: 84.11      1st Qu.: 515.3      1st Qu.:0.11660      1st Qu.:0.14720
## Median : 97.66      Median : 686.5      Median :0.13130      Median :0.21190
## Mean   :107.26      Mean    : 880.6      Mean    :0.13237      Mean    :0.25427
## 3rd Qu.:125.40      3rd Qu.:1084.0      3rd Qu.:0.14600      3rd Qu.:0.33910
## Max.   :251.20      Max.    :4254.0      Max.    :0.22260      Max.    :1.05800
## concavity_worst      concave.points_worst      symmetry_worst      fractal_dimension_worst
## Min.   :0.0000      Min.    :0.00000      Min.    :0.1565      Min.    :0.05504
## 1st Qu.:0.1145      1st Qu.:0.06493      1st Qu.:0.2504      1st Qu.:0.07146
## Median :0.2267      Median :0.09993      Median :0.2822      Median :0.08004
## Mean   :0.2722      Mean    :0.11461      Mean    :0.2901      Mean    :0.08395
## 3rd Qu.:0.3829      3rd Qu.:0.16140      3rd Qu.:0.3179      3rd Qu.:0.09208
## Max.   :1.2520      Max.    :0.29100      Max.    :0.6638      Max.    :0.20750
##
##      x
## Mode:logical
## NA's:569
##
##
##
##
```

The Breast Cancer Dataset contains a total number of 33 variables and 569 observations. The variables `x` and `id` are not interesting for our analysis, so we won't consider them later. There are no missing values in this data set, we don't need to perform imputation. All the variables are numeric, except `diagnosis`, which is stored as a logical variable, this must be taken into account for future steps.

Check diagnosis:

```
cat(" Number cases with malignant tumor: ",sum(breast_dataset$diagnosis == "M"),"\n", "Number cases with benign tumor: ",sum(breast_dataset$diagnosis == "B"))
```

```
## Number cases with malignant tumor: 212
## Number cases with benign tumor: 357
```

```
cat(" From a total number of", nrow(breast_dataset), "tumor cases, ", (sum(breast_dataset$diagnosis == "M")/nrow(breast_data
set))*100, "% are malignant tumors.")
```


```
## From a total number of 569 tumor cases, 37.25835 % are malignant tumors.
```

Pearson correlation:

```
# Select part of the original data, as some variables are not of our interest:
breast_df <- breast_dataset[,3:ncol(breast_dataset)-1]
breast_df$diagnosis <- as.integer(factor(breast_df$diagnosis))-1

breast_correlations <- cor(breast_df,method="pearson")

# Generate correlation plot:
corrplot(breast_correlations, hclust.method = "ward",method = "square",
         order = "FPC", type = "full", tl.col = "darkslategray",tl.cex=0.5)
```



As it is a huge matrix it is better not to include in the report the table of correlation values.

Note that diagnosis has been changed, now Benign tumors are indicated as 0 and Malignant tumors as 1.

MODEL BUILDING

Before starting with model building, we must pre-process the data. We will change the type of variable diagnosis and normalize data (with the custom made function *normalize*):

```
# Create function:
normalize <- function(x) {
  return ((x - min(x)) / (max(x) - min(x)))
}

# Generate a seed to replicate later the results :
set.seed(123)

# Diagnosis as an integer:
breast_dataset$diagnosis = as.integer(factor(breast_dataset$diagnosis))-1

# Normalize data:
breast_dataset_norm <- as.data.frame(lapply(breast_dataset[2:31], normalize))
```

Now we move to the model building. The first step is to separate the data in training and test:

```
# Generate the split object:
breast_split <- data_split <- initial_split(breast_dataset_norm, prop = 3/4)

# Build the training breast data set (with 3/4 of the data)
breast_training <- breast_split %>% training()
# Obtain labels, it's necessary to keep track of labels but in a separate variable:
breast_training_labels <- breast_training[, 1]
# Remove labels from training set:
breast_training <- as.data.frame(breast_training[-1])

# Build the testing breast data set:
breast_test <- breast_split %>% testing()
# Obtain labels:
breast_test_labels <- breast_test[, 1]
# Remove labels from test set:
breast_test <- as.data.frame(breast_test[-1])
```

KNN

In this part of the report, we will build and test a model based on **K-nearest neighbors (KNN)** algorithm.

KNN is a classification method that estimates the likelihood that a data point will become a member of one group or another, calculating Euclidean distances.

It is a supervised machine learning algorithm, we must provide the labels of the samples, in this case the diagnosis (Malignant (M), Benign (B)).

As an initial parameter we must choose the number of nearest neighbors to include (K). In order to generate a model with high accuracy, we must test different K values. In this report, we have K = 21 as the performance of the model was appropriate (high accuracy and low number of false positive and negative), but other K values that lead to lower accuracy models were tried previously. It was a trial-error process.

```
# Run KNN
set.seed(111)
knn_pred <- knn(train = breast_training, test = breast_test, cl = breast_training_labels, k=21,prob=TRUE)
```

Check the accuracy:

```
cat("The accuracy of the KNN model is:",sum(knn_pred == breast_test_labels)/nrow(breast_test))
```

```
## The accuracy of the KNN model is: 0.9370629
```

Check results:

```
CrossTable(x = breast_test_labels, y = knn_pred, prop.chisq=FALSE)
```

```
##
##
## Cell Contents
## |-----|
## |               N |
## |      N / Row Total |
## |      N / Col Total |
## |      N / Table Total |
## |-----|
##
##
## Total Observations in Table: 143
##
##
##
##      | knn_pred
## breast_test_labels |      0 |      1 | Row Total |
## -----|-----|-----|
##      0 |      79 |      1 |      80 |
##      | 0.988 | 0.012 | 0.559 |
##      | 0.908 | 0.018 |      |
##      | 0.552 | 0.007 |      |
## -----|-----|-----|
##      1 |      8 |     55 |     63 |
##      | 0.127 | 0.873 | 0.441 |
##      | 0.092 | 0.982 |      |
##      | 0.056 | 0.385 |      |
## -----|-----|-----|
##      Column Total |     87 |     56 |     143 |
##      | 0.608 | 0.392 |      |
## -----|-----|-----|
##
##
```


```
# Generate a Confusion Matrix:
conf_matrix <- data.frame(knn_pred,breast_test_labels)

# Rename Columns
names(conf_matrix) <- c("Predicted", "Actual")

conf_matrix <- conf_mat(conf_matrix, Actual,Predicted)

autoplot(conf_matrix, type = "heatmap") +
  scale_fill_gradient(low="slategray2",high = "mediumpurple3") +
  theme(legend.position = "right") + labs(title = "KNN Confusion Matrix")+
  theme(plot.title = element_text(hjust = 0.5))
```

```
## Scale for 'fill' is already present. Adding another scale for 'fill', which
## will replace the existing scale.
```



ROC CURVE:


```
knn.ROC <- roc(predictor=breast_test_labels, response=attributes(knn_pred)$prob)
```

```
## Warning in roc.default(predictor = breast_test_labels, response =
## attributes(knn_pred)$prob): 'response' has more than two levels. Consider
## setting 'levels' explicitly or using 'multiclass.roc' instead
```

```
## Setting levels: control = 0.571428571428571, case = 0.619047619047619
```

```
## Setting direction: controls < cases
```

```
plot(knn.ROC)
```



Area Under the Curve (AUC):

```
knn.ROC$auc
```

```
## Area under the curve: 0.75
```

The previous Cross Table and confusion matrix allows to identify True positives (55 cases with malignant tumor), True negatives (79 cases with benignant tumor).

We can also identify False negatives (8 cases that **KNN** predicted as negative, but in fact were positive) and False positives (1 case predicted as positive when it was negative).

The ROC plot and Area Under the curve gives us insights in the model performance. The closer the AUC value to 1, the better prediction. The **KNN** model implemented in the report has an AUC = 0.75, further changes should be done in order to improve it. For example, the model could be improved with a previous pre-processing of data doing a z-score standardization, instead of the normalization done in this report.

Decision Tree

Decision Tree is a classification algorithm that has a flowchart similar to a tree structure.

This model classifies instances by sorting them down the tree from the root to any leaf node, which provides the classification of the instance. It contains internal nodes that denote a test on a specific attribute, while each branch represents an outcome of the test. The terminal nodes contain class labels.

The instances are classified by sorting them from the root to a leaf node, that will provide the classification of the instance.

It is a non-parametric supervised learning method.

In order to generate the Decision Tree model we will use the function C5.0:

```
# Transform Labels to factor, necessary for the correct flow of the function:

breast_training$diagnosis<-as.factor(breast_training_labels)

breast_test['diagnosis'] = breast_test_labels
breast_test$diagnosis<-as.factor(breast_test$diagnosis)

# Set seed:
set.seed(1111)
# Generate the model:
dt_model <-C5.0(diagnosis ~., data = breast_training)

# Check:
dt_model
```

```
##
## Call:
## C5.0.formula(formula = diagnosis ~ ., data = breast_training)
##
## Classification Tree
## Number of samples: 426
## Number of predictors: 29
##
## Tree size: 7
##
## Non-standard options: attempt to group attributes
```

Now make predictions with out testing dataset:

```
dt_pred<- predict(dt_model, breast_test, type="class")
```

Check the accuracy:

```
cat("The accuracy of the Decision Tree model is:",sum(dt_pred == breast_test$diagnosis)/nrow(breast_test))
```

```
## The accuracy of the Decision Tree model is: 0.9300699
```

Now we have to compare the predicted labels to the real labels of the test dataset:

```
# Check results
CrossTable(x = breast_test_labels, y = dt_pred, prop.chisq=FALSE)
```

```
##
##
## Cell Contents
## |-----|
## |               N |
## | N / Row Total |
## | N / Col Total |
## | N / Table Total |
## |-----|
##
##
## Total Observations in Table: 143
##
##
##
## dt_pred
## breast_test_labels | 0 | 1 | Row Total |
## -----|-----|-----|
## 0 | 73 | 7 | 80 |
## | 0.912 | 0.087 | 0.559 |
## | 0.961 | 0.104 | |
## | 0.510 | 0.049 | |
## -----|-----|-----|
## 1 | 3 | 60 | 63 |
## | 0.048 | 0.952 | 0.441 |
## | 0.039 | 0.896 | |
## | 0.021 | 0.420 | |
## -----|-----|-----|
## Column Total | 76 | 67 | 143 |
## | 0.531 | 0.469 | |
## -----|-----|-----|
##
##
```


```
# Generate a Confusion Matrix:
conf_matrix2 <- data.frame(dt_pred,breast_test_labels)

# Rename Columns
names(conf_matrix2) <- c("Predicted", "Actual")

conf_matrix2 <- conf_mat(conf_matrix2, Actual,Predicted)

autoplot(conf_matrix2, type = "heatmap") +
  scale_fill_gradient(low="slategray2",high = "mediumpurple3") +
  theme(legend.position = "right") + labs(title = "Decision Tree Confusion Matrix")+
  theme(plot.title = element_text(hjust = 0.5))
```

```
## Scale for 'fill' is already present. Adding another scale for 'fill', which
## will replace the existing scale.
```



Comparing the Confusion matrices, for the **Decision Tree model**, we can see that the algorithm has a better performance than **KNN** when detecting positive cases, those that are malignant tumors. From a total number of 63 malignant tumors, it predicted 60. This result is better than KNN, which only predicted 55 cases. But on the other hand, we can see that with Decision Tree has a greater number of false positives (7 vs 1 in KNN), this model is prone to type 1 errors.

Neural Network

Neural Networks consist of an artificial network of functions. These functions also called parameters (or neurons), allow to learn and to modify internal elements by learning from new data. Each neuron after getting an input, produces an output. The outputs are sent to the next of neurons, and so on. The process continues all layers have been considered, and then the terminal neurons then output the final result of the model.

Generate a Neural Network model. The parameters are the ones recommended in literature:

```
set.seed(1111)
nn_model <- nnet(diagnosis ~., data=breast_training, size=15, rang = 1, decay = 8e-4, maxit = 100)
```

```
## # weights: 466
## initial value 257.930455
## iter 10 value 31.020161
## iter 20 value 14.462695
## iter 30 value 5.715691
## iter 40 value 4.426506
## iter 50 value 3.681904
## iter 60 value 3.274219
## iter 70 value 3.084947
## iter 80 value 2.982445
## iter 90 value 2.934833
## iter 100 value 2.902831
## final value 2.902831
## stopped after 100 iterations
```

Prediction using the previously generated Neural Network model:

```
nn_pred <- predict(nn_model, breast_test,type = c("class"))
```

Check the accuracy:

```
cat("The accuracy of the Neural Network model is:",sum(nn_pred == breast_test$diagnosis)/nrow(breast_test))
```

```
## The accuracy of the Neural Network model is: 0.972028
```

Check the numeric results in a table format:

```
CrossTable(breast_test$diagnosis, nn_pred, prop.chisq = FALSE,prop.c = FALSE, prop.r = FALSE, dnn = c("Actual", "Predicted"
))
```

```
##
##
##   Cell Contents
## |-----|
## |               N |
## | N / Table Total |
## |-----|
##
##
## Total Observations in Table: 143
##
##
##               | Predicted
##   Actual |      0 |      1 | Row Total |
## -----|-----|-----|-----|
##      0 |      78 |      2 |      80 |
##      | 0.545 | 0.014 |      |
## -----|-----|-----|-----|
##      1 |      2 |     61 |     63 |
##      | 0.014 | 0.427 |      |
## -----|-----|-----|-----|
## Column Total |      80 |     63 |     143 |
## -----|-----|-----|-----|
##
##
```


```
# Generate a Confusion Matrix:
conf_matrix3 <- data.frame(nn_pred,breast_test_labels)

# Rename Columns
names(conf_matrix3) <- c("Predicted", "Actual")

conf_matrix3 <- conf_mat(conf_matrix3, Actual,Predicted)


autoplot(conf_matrix3, type = "heatmap") +
  scale_fill_gradient(low="slategray2",high = "mediumpurple3") +
  theme(legend.position = "right") + labs(title = "Neural Network Confusion Matrix")+
  theme(plot.title = element_text(hjust = 0.5))
```

```
## Scale for 'fill' is already present. Adding another scale for 'fill', which
## will replace the existing scale.
```



Now plot the Neural-network as a diagram:

```
plotnet(nn_model, alpha = 0.6, circle_cex = 2, cex_val = 0.8, circle_col = "mediumorchid2", bord_col = "mediumorchid2", prune_col = TRUE)
```



If we compare with the results of **Neural Network** with the previous Confusion Matrices of **KNN** and **Decision Tree**, we can see that this model has the best performance. It has the lowest number of false positives as well as false negatives. Also, we can see that the accuracy is very high.

We can also check the diagram and see the conformation of layers of the model. ## Logistic Regression

Logistic Regression is another classification algorithm that can be used to assign observations to a discrete set.

It is based on probabilities and the cost function is Sigmoid (or Logistic Function).

```
lr_model <- train(diagnosis~., data=breast_training, method="glm", family=binomial())  
varImp(lr_model)
```

```
## glm variable importance
##
##    only 20 most important variables shown (out of 29)
##
##                                Overall
## concavity_se                100.00
## compactness_mean           94.26
## smoothness_mean            84.20
## fractal_dimension_mean      73.76
## concave.points_se           59.38
## radius_mean                 58.85
## radius_worst                 53.18
## area_worst                   46.93
## smoothness_se               41.62
## symmetry_worst               41.52
## fractal_dimension_se        40.20
## perimeter_mean              36.58
## symmetry_mean                35.86
## symmetry_se                  34.51
## compactness_se               32.97
## perimeter_se                 30.89
## concavity_worst              24.88
## area_mean                    24.52
## concave.points_mean          15.42
## radius_se                    13.20
```

Only 20 most important variables shown (out of 29).

Prediction:

```
set.seed(11111)
lr_pred<-predict(lr_model,breast_test)
```

Check the accuracy:

```
cat("The accuracy of the Logistic Regression is:",sum(lr_pred == breast_test$diagnosis)/nrow(breast_test))
```

```
## The accuracy of the Logistic Regression is: 0.958042
```

Check the numeric results in a table format:

```
CrossTable(x = breast_test_labels, y = lr_pred, prop.chisq=FALSE)
```

```
##
##
## Cell Contents
## |-----|
## |               N |
## | N / Row Total |
## | N / Col Total |
## | N / Table Total |
## |-----|
##
##
## Total Observations in Table: 143
##
##
##
## lr_pred
## breast_test_labels | 0 | 1 | Row Total |
## -----|-----|-----|
## 0 | 75 | 5 | 80 |
## | 0.938 | 0.062 | 0.559 |
## | 0.987 | 0.075 | |
## | 0.524 | 0.035 | |
## -----|-----|-----|
## 1 | 1 | 62 | 63 |
## | 0.016 | 0.984 | 0.441 |
## | 0.013 | 0.925 | |
## | 0.007 | 0.434 | |
## -----|-----|-----|
## Column Total | 76 | 67 | 143 |
## | 0.531 | 0.469 | |
## -----|-----|-----|
##
##
```

```
# Generate a Confusion Matrix:
conf_matrix4 <- data.frame(lr_pred,breast_test_labels)

# Rename Columns
names(conf_matrix4) <- c("Predicted", "Actual")

conf_matrix4 <- conf_mat(conf_matrix4, Actual,Predicted)

autoplot(conf_matrix4, type = "heatmap") +
  scale_fill_gradient(low="slategray2",high = "mediumpurple3") +
  theme(legend.position = "right") + labs(title = "Logistic Regression Confusion Matrix")+
  theme(plot.title = element_text(hjust = 0.5))
```

```
## Scale for 'fill' is already present. Adding another scale for 'fill', which
## will replace the existing scale.
```



From the Confusion Matrix we can see that this model has a greater number of false positives than the Neural Network, but a lower number of false negatives (we also have to take into account that the difference is not that big, only by 1 case).

This model is better at not mislabeling malignant tumors as benign, but instead, it mislabels a greater number of benign tumors as malignant.

Which of all models performs better for this data? Discuss.

According to the accuracy, the rank would be (from higher accuracy to lower): Neural Network, Logistic Regression, KNN and Decision Tree.

Of all the models, the **Neural Network** model is the one with a better performance (it has the highest accuracy value, 0.972028).